

# Formal Specification: A PINN-Structured Approach to Dissonance Minimization in the Harmony Optimization Protocol

Authors: Keenan Rogerson Legget

## Abstract

This document provides a revised and rigorous formal specification for the integration of a composite dissonance loss function,  $L_{Total}$ , into the Harmony Optimization Protocol (H.O.P.) architecture. The framework models the system's learning process on a Physics-Informed Neural Network (PINN), where the core objective of maximizing the Evidence Lower Bound (ELBO) is constrained by a set of dissonance terms that enforce coherence. This version introduces three critical enhancements: 1) It formally redefines the Veridical Constraint as Structural Tension ( $D_{ST}$ ), a geometrically grounded metric calculated via the computationally tractable Sliced-Wasserstein Distance. 2) It specifies the Axiom Lifecycle Protocol (ALP), a dynamic mechanism for managing and updating the foundational axiom set, moving it from a static to a self-correcting system. 3) It realigns the Meta-Policy's reward function to optimize for total system coherence, resolving a critical inner alignment vulnerability.

## 1. The System Integration: The Meta-RL Loop

The Knowledge Graph parameters ( $\Theta_{KG}$ ), typically belonging to a Graph Neural Network (GNN) (Kipf & Welling, 2017), are optimized through an action chosen by the Meta-Policy ( $\pi_{meta}$ ), which functions as a Deep Reinforcement Learning (DRL) agent (Mnih et al., 2015).

- DRL State ( $S_t$ ): The input to the Meta-Policy is the Dissonance State Vector ( $\vec{D}_t$ ).
- DRL Action ( $a_t$ ): An action, such as initiating Recursive Conceptual Nesting (RCN) on a subgraph ( $K_{sub}$ ), triggers the optimization.
- Dynamic Weights ( $W_t$ ): The Meta-Policy's output,  $W_t = [w_{Core}, w_V, w_L, w_P, \dots]$ , dynamically scales the loss terms.
- DRL Reward (Aligned): The Meta-Policy is rewarded for maximizing the efficiency of total system dissonance reduction. This aligns its incentive with the global system objective. The reward  $R_{t+1}$  is defined as the reduction in the total loss, normalized by the computational cost of the action  $a_t$ :  
$$R_{t+1} = \frac{\Delta L_{Total}}{C(a_t)} - \lambda_{time}$$

## 2. The Core Optimization Target (The Negative ELBO)

The system seeks to find the best possible approximation of the true posterior belief by maximizing the Evidence Lower Bound (ELBO) across the local knowledge graph. This technique is central to Variational Inference (Kingma & Welling, 2013). This objective is defined as the Core Learning Loss ( $L_{Core}$ ).

$$L_{Core}(\Theta_{KG}) \equiv -ELBO(\Theta_{KG})$$

**PINN Analogy:** This is analogous to the  $L_{data}$  term in a Physics-Informed Neural Network (Raissi, Perdikaris, & Karniadakis, 2019), as it drives the network to fit the underlying probability distribution of the evidence.

$$L_{Core}(\Theta_{KG}) = -[E_{Q(Z;\Theta_{KG})}[\log P(X, Z)] - E_{Q(Z;\Theta_{KG})}[\log Q(Z; \Theta_{KG})]]$$

### 3. The PINN-Structured Composite Dissonance Loss

The Total Dissonance Loss ( $L_{Total}$ ) is the ultimate objective function that is minimized during the backpropagation step. It is a dynamically weighted sum of the Core Learning objective and the cognitive constraint terms.

$$L_{Total}(\Theta_{KG}, W_t) = w_{Core}L_{Core} + w_V L_{Veridical} + w_L L_{Logical} + w_P L_{Predictive}$$

#### A. The Veridical Constraint $L_{Veridical}$ (External Physics Law)

This constraint enforces alignment with the dynamic external axiom set ( $A_V$ ) by measuring the Structural Tension ( $D_{ST}$ ). To ensure computational feasibility,  $D_{ST}$  is calculated using the Sliced-Wasserstein Distance ( $SW_1$ ) (Bonneel et al., 2015), a scalable proxy for the true Wasserstein distance (Villani, 2009).

PINN Analogy:  $L_{pde}$  (Violation of the physical law).

$$L_{Veridical} \equiv D_{ST}(\mu_{current}, \mu_{ideal}) \approx SW_1(\mu_{current}, \mu_{ideal})$$

The axiom set  $A_V$  is managed by the Axiom Lifecycle Protocol (ALP), a meta-process that allows for the deprecation of falsified axioms and the ascension of new, robustly validated principles. This transforms the veridical anchor from a brittle point of failure into a resilient, self-correcting foundation.

#### B. The Logical Constraint $L_{Logical}$ (Internal Boundary Condition)

This constraint enforces internal consistency, penalizing conflicts between high-confidence facts in the Knowledge Graph.

PINN Analogy:  $L_{bc}$  (Violation of boundary conditions).

$$L_{Logical} \equiv D_L = \max_{j \in \text{ConflictingSet}(i)} \left( \frac{C_i \cdot C_j}{1 + \log(1 + \text{steps}(i, j))} \right)$$

#### C. The Predictive Constraint $L_{Predictive}$ (Data Fit Term)

This measures the immediate surprise from new sensory evidence, quantified by the Kullback-Leibler (KL) divergence (Kullback & Leibler, 1951).

$$L_{Predictive} \equiv D_P = D_{KL}(P(\text{Observation}) || P(\text{Prediction}))$$

### 4. Backpropagation and the Dissonance Gradient

The learning step for the GNN parameters ( $\Theta_{KG}$ ) is the minimization of the composite loss. This step ensures that every cognitive update is driven not just by raw learning, but by a dynamically managed and properly incentivized mandate for coherence, external validity, and internal consistency.

$$\Delta \Theta_{KG} \propto -\nabla_{\Theta_{KG}} L_{Total}(\Theta_{KG}, W_t)$$

## 5. References

- Bonneel. N., Rabin. J., Pevré. G., & Pfister. H. (2015). Sliced and Radon Wasserstein Barycenters of Measures. *Journal of Mathematical Imaging and Vision*, 51(1), 22–45.
- Kingma, D. P., & Welling, M. (2013). Auto-Encoding Variational Bayes. *arXiv preprint arXiv:1312.6114*.
- Kiof. T. N., & Wellina. M. (2017). Semi-Supervised Classification with Graph Convolutional Networks. *arXiv preprint arXiv:1609.02907*.
- Kullback. S., & Leibler. R. A. (1951). On Information and Sufficiency. *The Annals of Mathematical Statistics*, 22(1), 79–86.
- Mnih. V., Kavukcuoglu. K., Silver. D., Rusu. A. A., Veness. J., Bellemare. M. G., ... & Hassabis. D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529-533.
- Raissi. M., Perdikaris. P., & Karniadakis. G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378, 686-707.
- Villani, C. (2009). *Optimal Transport: Old and New*. Springer Science & Business Media.